

SICP Exercise Solutions for Section 1.2

Paul L. Snyder

March 24, 2014

Contents

1	Section 1.2.1 Linear Recursion and Iteration	3
1.1	Exercise 1.9: Thinking about procedures and processes	3
1.1.1	Problem	3
1.1.2	Answer	3
1.2	Exercise 1.10: Ackermann's function	5
2	Section 1.2.2 Tree Recursion	7
2.1	Exercise 1.11: Converting a recursive process to an iterative process	7
2.1.1	Problem	7
2.1.2	Answer	7
2.2	Exercise 1.12: Computing Pascal's triangle	9
2.2.1	Problem	9
2.2.2	Answer	10
2.3	Exercise 1.13: A Fibonacci proof	11
2.3.1	Problem	11
2.3.2	Answer	11
3	Section 1.2.3 Orders of Growth	14
3.1	TODO Exercise 1.14: Counting change	14
3.1.1	Problem	14
3.1.2	Answer	14
3.2	Exercise 1.15: Approximating sine	16
3.2.1	a.	17
3.2.2	b.	17

4	Section 1.2.4 Exponentiation	18
4.1	Exercise 1.16: Iterative exponentiation in logarithmic time . .	18
4.1.1	Problem	18
4.1.2	Answer	18
4.2	Exercise 1.17: Recursive integer multiplication with square and halve	19
4.2.1	Problem	19
4.2.2	Answer	20
4.3	Exercise 1.18: Iterative integer multiplication with square and halve	20
4.3.1	Problem	20
4.3.2	Answer	21
4.4	Exercise 1.19: Logarithmic Fibonacci calculations	21
4.4.1	Problem	21
4.4.2	Answer	22
5	Section 1.2.5 Greatest Common Divisors	24
5.1	TODO Exercise 1.20: Revisiting applicative order and normal order	24
5.1.1	Problem	24
5.1.2	Answer	24
6	Section 1.2.6 Testing for Primality	24
6.1	Exercise 1.21: Using smallest-divisor	24
6.1.1	Problem	24
6.1.2	Answer	24
6.2	Exercise 1.22: Measuring runtime	25
6.2.1	Problem	25
6.2.2	Answer	26
6.3	Exercise 1.23: Speeding up smallest-divisor	29
6.3.1	Problem	29
6.3.2	Answer	29
6.4	Exercise 1.24: Putting the Fermat method to work	31
6.4.1	Problem	31
6.4.2	Answer	31
6.5	Exercise 1.25: A not-so-fast use of fast-expt	33
6.5.1	Problem	33
6.5.2	Answer	33
6.6	Exercise 1.26: A subtle slowdown in expmod	36
6.6.1	Answer	36

6.7	Exercise 1.27: Fooling Fermat with Carmichael numbers . . .	37
6.7.1	Problem	37
6.7.2	Answer	37
6.8	Exercise 1.28: The Miller-Rabin test	39
6.8.1	Problem	39
6.8.2	Answer	39

1 Section 1.2.1 Linear Recursion and Iteration

1.1 Exercise 1.9: Thinking about procedures and processes

1.1.1 Problem

Each of the following two procedures defines a method for adding two positive integers in terms of the procedures `inc`, which increments its argument by 1, and `dec`, which decrements its argument by 1.

```
(define (+ a b)
  (if (= a 0)
      b
      (inc (+ (dec a) b))))
```

```
(define (+ a b)
  (if (= a 0)
      b
      (+ (dec a) (inc b))))
```

Using the substitution model, illustrate the process generated by each procedure in evaluating `(+ 4 5)`. Are these processes iterative or recursive?

1.1.2 Answer

First procedure:

```
(+ 4 5)
(if (= 4 0)
    5
    (inc (+ (dec 4) 5)))
(inc (+ (dec 4) 5))
(inc (+ 3 5))
(inc (if (= 3 5)
```

```

5
    (inc (+ (dec 3) 5)))
(inc (inc (+ (dec 3) 5)))
(inc (inc (+ 2 5) ))
(inc (inc (if (= 2 5)
5
    (inc (+ (dec 2) 5)))))
(inc (inc (inc (+ (dec 2) 5))))
(inc (inc (inc (+ 1 5))))
(inc (inc (inc (if (= 1 0)
5
    (inc (+ (dec 1) 5)))))
(inc (inc (inc (inc (+ (dec 1) 5)))))
(inc (inc (inc (inc (+ 0 5)))))
(inc (inc (inc (inc (if (= 0 0)
5
    (inc (+ (dec 0) 5)))))
(inc (inc (inc (inc 5))))
(inc (inc (inc 6)))
(inc (inc 7))
(inc 8)
9

```

Second procedure:

```

(+ 4 5)
(if (= 4 0)
5
    (+ (dec 4) (inc 5)))
(+ (dec 4) (inc 5))
(+ 3 6)
(if (= 3 0)
6
    (+ (dec 3) (inc 6)))
(+ (dec 3) (inc 6))
(+ 2 7)
(if (= 2 0)
7
    (+ (dec 2) (inc 7)))
(+ (dec 2) (inc 7))

```

```

(+ 1 8)
(if (= 1 0)
    8
    (+ (dec 1) (inc 8)))
(+ (dec 1) (inc 8))
(+ 0 9)
(if (= 0 0)
    9
    (+ (dec 0) (inc 9)))
9

```

As can be seen, the first procedure is a linear recursive process, building up deferred operations. The second is a linear iterative process.

1.2 Exercise 1.10: Ackermann's function

The following procedure computes a mathematical function called Ackermann's function.

```

(define (A x y)
  (cond ((= y 0) 0)
        ((= x 0) (* 2 y))
        ((= y 1) 2)
        (else (A (- x 1)
                   (A x (- y 1))))))

```

What are the values of the following expressions?

(A 1 10)

1024

(A 2 4)

65536

(A 3 3)

65536

Consider the following procedures, where **A** is the procedure defined above:

```

(define (f n) (A 0 n))

(define (g n) (A 1 n))

(define (h n) (A 2 n))

(define (k n) (* 5 n n))

(define (investigate name func num)

  (define (investigate-iter i)
    (cond ((< i num)
           (printf " ~a:~a" i (func i))
           (investigate-iter (+ i 1)))
          (else (newline))))
  (display name)
  (investigate-iter 0))

(investigate "k" k 5)

```

Give concise mathematical definitions for the functions computed by the procedures `f`, `g`, and `h` for positive integer values of n . For example, `(k n)` computes $5n^2$.

```
(investigate "f" f 17)
```

The output of `f` is straightforward to analyze:

$$f(n) => 2 * n$$

```
(investigate "g" g 17)
```

Similarly, `g` is easy to identify for anyone who's been programming for any length of time:

$$g(n) => 2^n$$

```
(investigate "h" h 5)
```

`h`, on the other hand, is a serious pain!

(A 2 1) => 2	(2^(2^0))
(A 2 2) => (A 1 (A 2 1)) => (A 1 2) => 4	(2^(2^1)) (2^(2^(2^0)))
(A 2 3) => (A 1 (A 2 2)) => (A 1 4) => 16	(2^(2^2)) (2^(2^(2^1)))
(A 2 4) => (A 1 (A 2 3)) => (A 1 16) => 65536	(2^(2^4)) (2^(2^(2^2)))
(A 2 5) => (A 1 (A 2 4)) => (A 1 65536) => 20035...56736	(2^(2^16)) (2^(2^(2^4)))

$2^2 \wedge 2^1$
 $2^3 \wedge 2^2 \wedge 2^1$

$h(1) \Rightarrow 2$
 $h(2) \Rightarrow 2 \wedge 2$
 $h(3) \Rightarrow 2 \wedge 2 \wedge 2$
 $h(4) \Rightarrow 2 \wedge 2 \wedge 2 \wedge 2$
 ...

There's no standard mathematical notation for this pattern; which is commonly called a "power tower". A more formal term is *tetration*.

2 Section 1.2.2 Tree Recursion

2.1 Exercise 1.11: Converting a recursive process to an iterative process

2.1.1 Problem

A function f is defined by the rule that $f(n) = n$ if $n < 3$ and $f(n) = f(n-1) + 2f(n-2) + 3f(n-3)$ if $n \geq 3$. Write a procedure that computes f by means of a recursive process. Write a procedure that computes f by means of an iterative process.

2.1.2 Answer

```

(define (f-rec n)
  (if (< n 3)
      n
      (+ (f-rec (- n 1))
         (* 2 (f-rec (- n 2)))
         (* 3 (f-rec (- n 3))))))

(define (f-iter n)

```

```

(define (f-inner i fn1 fn2 fn3)
  (define cur (+ fn1 (* 2 fn2) (* 3 fn3)))
  (if (= i n)
      cur
      (f-inner (+ i 1) cur fn1 fn2)))

(if (< n 3)
    n
    (f-inner 3 2 1 0)))

(define (compare f1 f2 k)
  (define (compare-iter i good?)
    (define t0 (current-milliseconds))
    (define r1 (f1 i))
    (define t1 (current-milliseconds))
    (define r2 (f2 i))
    (define t2 (current-milliseconds))

    (printf "~a: ~a (~a ms) ~a (~a ms) => ~a~n"
            i
            r1 (- t1 t0)
            r2 (- t2 t1)
            (= r1 r2))
    (if (< i k)
        (compare-iter (+ i 1) (and good? (= r1 r2)))
        (and good? (= r1 r2))))

  (compare-iter 1 true))

(compare f-rec f-iter 33)

1: 1 (0 ms) 1 (0 ms) => #t
2: 2 (0 ms) 2 (0 ms) => #t
3: 4 (0 ms) 4 (0 ms) => #t
4: 11 (0 ms) 11 (0 ms) => #t
5: 25 (0 ms) 25 (0 ms) => #t
6: 59 (0 ms) 59 (0 ms) => #t
7: 142 (0 ms) 142 (0 ms) => #t
8: 335 (0 ms) 335 (0 ms) => #t
9: 796 (0 ms) 796 (0 ms) => #t

```



```

10: 1892 (0 ms) 1892 (0 ms) => #t
11: 4489 (0 ms) 4489 (0 ms) => #t
12: 10661 (0 ms) 10661 (0 ms) => #t
13: 25315 (0 ms) 25315 (0 ms) => #t
14: 60104 (0 ms) 60104 (0 ms) => #t
15: 142717 (0 ms) 142717 (0 ms) => #t
16: 338870 (1 ms) 338870 (0 ms) => #t
17: 804616 (0 ms) 804616 (0 ms) => #t
18: 1910507 (1 ms) 1910507 (0 ms) => #t
19: 4536349 (1 ms) 4536349 (0 ms) => #t
20: 10771211 (2 ms) 10771211 (0 ms) => #t
21: 25575430 (4 ms) 25575430 (0 ms) => #t
22: 60726899 (8 ms) 60726899 (0 ms) => #t
23: 144191392 (13 ms) 144191392 (0 ms) => #t
24: 342371480 (26 ms) 342371480 (0 ms) => #t
25: 812934961 (46 ms) 812934961 (0 ms) => #t
26: 1930252097 (84 ms) 1930252097 (0 ms) => #t
27: 4583236459 (154 ms) 4583236459 (0 ms) => #t
28: 10882545536 (288 ms) 10882545536 (0 ms) => #t
29: 25839774745 (522 ms) 25839774745 (0 ms) => #t
30: 61354575194 (963 ms) 61354575194 (0 ms) => #t
31: 145681761292 (1821 ms) 145681761292 (0 ms) => #t
32: 345910235915 (3296 ms) 345910235915 (0 ms) => #t
33: 821337484081 (6081 ms) 821337484081 (0 ms) => #t

```

2.2 Exercise 1.12: Computing Pascal's triangle

2.2.1 Problem

The following pattern of numbers is called "Pascal's triangle".

```

      1
     1 1
    1 2 1
   1 3 3 1
  1 4 6 4 1

```

The numbers at the edge of the triangle are all 1, and each number inside the triangle is the sum of the two numbers above it. Write a procedure that computes elements of Pascal's triangle by means of a recursive process.

2.2.2 Answer

```
;; Find the m-th number in the n-th row of Pascal's triangle
(define (pascal n m)
  (cond ((or (> m n) (< m 1) (< n 1)) -1) ; Error condition
        ((or (= m 1) (= m n)) 1) ; Outer numbers
        (else (+ (pascal (- n 1) (- m 1))
                  (pascal (- n 1) m)))))

;; Display first k rows of Pascal's triangle
(define (display-pascal k)
  (define (display-pascal-inner i j)
    (display (pascal i j))
    (cond ((< j i) (display " ") (display-pascal-inner i (+ j 1)))
          ((= i k) (newline))
          ((= i j) (newline) (display-pascal-inner (+ i 1) 1))))

  (display-pascal-inner 1 1))

(time (display-pascal 18))

1
1 1
1 2 1
1 3 3 1
1 4 6 4 1
1 5 10 10 5 1
1 6 15 20 15 6 1
1 7 21 35 35 21 7 1
1 8 28 56 70 56 28 8 1
1 9 36 84 126 126 84 36 9 1
1 10 45 120 210 252 210 120 45 10 1
1 11 55 165 330 462 462 330 165 55 11 1
1 12 66 220 495 792 924 792 495 220 66 12 1
1 13 78 286 715 1287 1716 1716 1287 715 286 78 13 1
1 14 91 364 1001 2002 3003 3432 3003 2002 1001 364 91 14 1
1 15 105 455 1365 3003 5005 6435 6435 5005 3003 1365 455 105 15 1
1 16 120 560 1820 4368 8008 11440 12870 11440 8008 4368 1820 560 120 16 1
1 17 136 680 2380 6188 12376 19448 24310 24310 19448 12376 6188 2380 680 136 17 1
cpu time: 16 real time: 15 gc time: 0
```

2.3 Exercise 1.13: A Fibonacci proof

2.3.1 Problem

Prove that $\text{Fib}(n)$ is the closest integer to $\phi^n/\sqrt{5}$, where $\phi = (1 + \sqrt{5})/2$. Hint: Let $\psi = (1 - \sqrt{5})/2$. Use induction and the definition of the Fibonacci numbers (see section 1.2.2) to prove that $\text{Fib}(n) = (\phi^n - \psi^n)/\sqrt{5}$.

2.3.2 Answer

This one is pretty hairy, and requires a modicum of mathematical sophistication to work through. For a proof like this (particularly when given a broad hint like this), the best approach is frequently to leap in and see where things take you.

In addition to the definition of ψ in the exercise, we also have a pointer to section 1.2.2, which gives us the definition of the Fibonacci sequence:

$$\text{Fib}(n) = \begin{cases} 0 & \text{if } n = 0 \\ 1 & \text{if } n = 1 \\ \text{Fib}(n-1) + \text{Fib}(n-2) & \text{otherwise} \end{cases}$$

Significantly, it also gives us a couple of facts about ϕ that will prove to be useful, as well:

$$\begin{aligned}\phi &= 1 + \sqrt{5}/2 \approx 1.6180 \\ \phi^2 &= \phi + 1\end{aligned}$$

Since this is a proof by induction, the first thing to do is to demonstrate the base cases. In this case, we have two, $n = 0$ and $n = 1$.

Starting with $n = 0$, we need to show $\text{Fib}(0) = 0$.

$$\text{Fib}(0) = \frac{\phi^0 - \psi^0}{\sqrt{5}} = \frac{1 - 1}{\sqrt{5}} = \frac{0}{\sqrt{5}} = 0$$

Easy! Now we show $\text{Fib}(1) = 1$.

$$\begin{aligned}
\text{Fib}(1) &= \frac{\phi^1 - \psi^1}{\sqrt{5}} \\
&= \frac{\frac{1+\sqrt{5}}{2} - \frac{1-\sqrt{5}}{2}}{\sqrt{5}} \\
&= \frac{\frac{1+\sqrt{5}-1+\sqrt{5}}{2}}{\sqrt{5}} \\
&= \frac{\frac{2\sqrt{5}}{2}}{\sqrt{5}} = \frac{\sqrt{5}}{\sqrt{5}} = 1
\end{aligned}$$

These are straightforward enough. Next, we need to assume that the hypothesis holds for $n - 1$ and n , and use this to show that this implies it also must hold for $n + 1$. That is, if we assume

$$\text{Fib}(n - 1) = \frac{\phi^{n-1} - \psi^{n-1}}{\sqrt{5}}$$

and

$$\text{Fib}(n) = \frac{\phi^n - \psi^n}{\sqrt{5}}$$

we show that it implies that

$$\text{Fib}(n + 1) = \frac{\phi^{n+1} - \psi^{n+1}}{\sqrt{5}}$$

Once again, the best thing to do is just leap in, start substituting, and see what results.

$$\begin{aligned}
\text{Fib}(n + 1) &= \text{Fib}(n - 1) + \text{Fib}(n) \\
&= \frac{\phi^{n-1} - \psi^{n-1}}{\sqrt{5}} + \frac{\phi^n - \psi^n}{\sqrt{5}} \\
&= \frac{\phi^{n-1} - \psi^{n-1} + \phi^n - \psi^n}{\sqrt{5}} \\
&= \frac{(\phi^n + \phi^{n-1}) - (\psi^n + \psi^{n-1})}{\sqrt{5}}
\end{aligned}$$

It's at this point where the useful tidbit about ϕ from the text comes in handy, that $\phi^2 = \phi + 1$. (As a note, it would have been helpful if the

exercise had mentioned that it is also the case that $\psi^2 = \psi + 1$!) With these identities in hand, we can factor out ϕ^{n-1} and ψ^{n-1} from the terms in the numerator and substitute accordingly.

$$\begin{aligned}
&= \frac{\phi^{n-1}(\phi + 1) - \psi^{n-1}(\psi + 1)}{\sqrt{5}} \\
&= \frac{\phi^{n-1}\phi^2 - \psi^{n-1}\psi^2}{\sqrt{5}} \\
&= \frac{\phi^{n+1} - \psi^{n+1}}{\sqrt{5}}
\end{aligned}$$

This concludes the inductive portion of the proof, but it doesn't get us all the way to what we were supposed to demonstrate: that, therefore, $\text{Fib}(n)$ is the closest integer to $\frac{\phi^n}{\sqrt{5}}$.

What would it mean for this to be the case? As the $\text{Fib}(n)$ will always be an integer, the distance between it and the other term must be less than $\frac{1}{2}$; if it were not the case then ϕ^n would be closer to a different integer. As an edge case, the difference can't be exactly $\frac{1}{2}$, as that would mean ϕ^n was equally close to a second integer. We can specify algebraically using the absolute value:

$$\left| \text{Fib}(n) - \frac{\phi^n}{\sqrt{5}} \right| < \frac{1}{2}$$

Trying out some substitutions, we can get this down to something manageable.

$$\left| \text{Fib}(n) - \frac{\phi^n}{\sqrt{5}} \right| = \left| \frac{\phi^n - \psi^n}{\sqrt{5}} - \frac{\phi^n}{\sqrt{5}} \right| = \left| -\frac{\psi^n}{\sqrt{5}} \right| = \left| \frac{\psi^n}{\sqrt{5}} \right|$$

This lets us examine what we are trying to prove in the light of this new equivalence:

$$\begin{aligned}
\left| \text{Fib}(n) - \frac{\phi^n}{\sqrt{5}} \right| &< \frac{1}{2} \\
\left| \frac{\psi^n}{\sqrt{5}} \right| &< \frac{1}{2} \\
|\psi^n| &< \frac{\sqrt{5}}{2}
\end{aligned}$$

We can see that

$$\frac{\sqrt{5}}{2} \approx 1.11803$$

and by definition,

$$\psi = \frac{1 - \sqrt{5}}{2} \approx -0.61803$$

Raising a real number $n \neq 0$ in the range $-1 < n < 1$ to a positive integer power can never result in a value that is outside that range. Thus, it can never exceed $\frac{\sqrt{5}}{2}$, since the latter is greater than one. This is sufficient to demonstrate that $\text{Fib}(n)$ is the closest integer to ϕ^n .

□

3 Section 1.2.3 Orders of Growth

3.1 TODO Exercise 1.14: Counting change

3.1.1 Problem

Draw the tree illustrating the process generated by the `count-change` procedure of section *Note 1-2-2:: in making change for 11 cents. What are the orders of growth of the space and number of steps used by this process as the amount to be changed increases?

3.1.2 Answer

First, we'll setup some tooling to output to GraphViz's `dot` format. This could be done in more complex (and interesting) ways, but this tries to stick as closely as possible to the Scheme features that have been discussed in the book so far. The additional features used are `format` and `printf` (for displaying output) and `random` for creating a sort-of-unique ID. If we were to just use the information available in a procedure (that is, its name and the parameters with which it was called), we wouldn't have a tree, as multiple calls to the same procedure with the same parameters would be collapsed.

```
(define (random-id)
  (random 5000000))
```

```
(define (make-name str)
  ;; Append random number to given string for a hopefully unique node
  ;; name. This isn't perfect, as there is a small possibility that
  ;; IDs could be repeated. As we aren't using assignment yet, this
```

```

;; is probably good enough.
(format "~a_~a" str (random-id)))

(define (dot-node name label)
  (printf "    ~a [shape=box,label=\"~a\"];~n" name label))

(define (dot-edge parent child)
  (printf "    ~a -> ~a;~n" parent child))

Now, the code itself with the above instrumentation included.

(define (count-change amount)
  (define name (make-name "count_change"))
  (dot-node name (format "(count-change ~a)" amount))
  (cc amount 5 name))

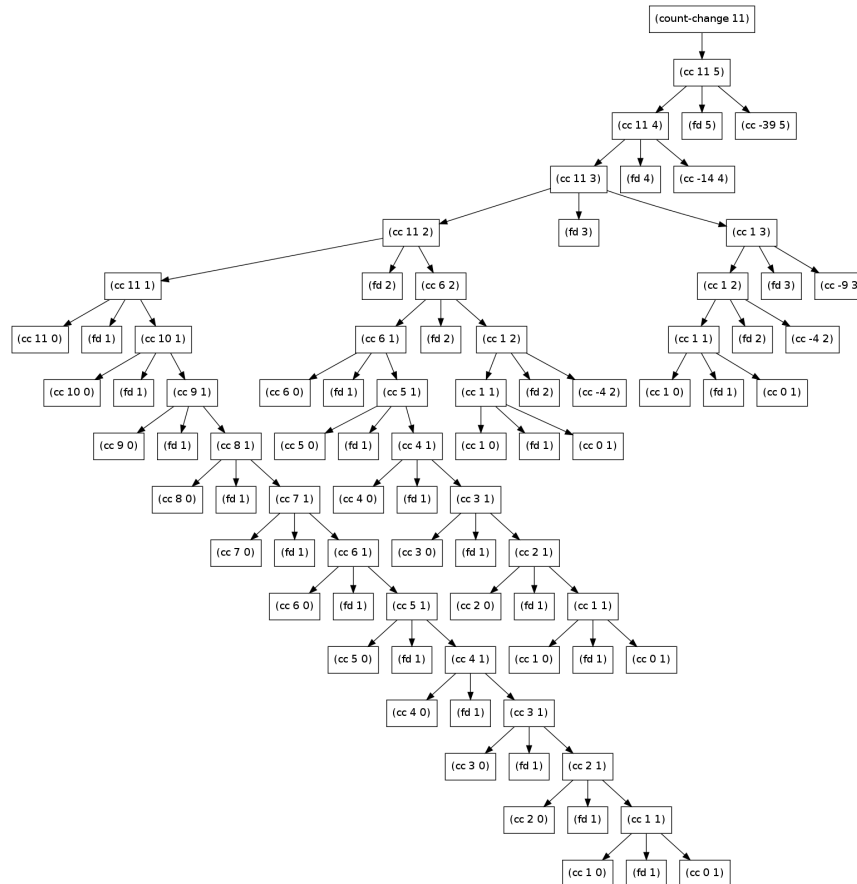
(define (cc amount kinds-of-coins parent)
  (define name (make-name "cc"))
  (dot-node name (format "(cc ~a ~a)" amount kinds-of-coins))
  (dot-edge parent name)
  (cond ((= amount 0) 1)
        ((or (< amount 0) (= kinds-of-coins 0)) 0)
        (else (+ (cc amount
                      (- kinds-of-coins 1)
                      name)
                  (cc (- amount
                        (first-denomination kinds-of-coins name))
                      kinds-of-coins
                      name))))))

(define (first-denomination kinds-of-coins parent)
  (define name (make-name "fd"))
  (dot-node name (format "(fd ~a)" kinds-of-coins))
  (dot-edge parent name)

  (cond ((= kinds-of-coins 1) 1)
        ((= kinds-of-coins 2) 5)
        ((= kinds-of-coins 3) 10)
        ((= kinds-of-coins 4) 25)
        ((= kinds-of-coins 5) 50)))

```

```
(count-change 11)
```



3.2 Exercise 1.15: Approximating sine

The sine of an angle (specified in radians) can be computed by making use of the approximation $\sin x \approx x$ if x is sufficiently small, and the trigonometric identity

$$\sin x = 3 \sin \frac{x}{3} - 4 \sin^3 \frac{x}{3}$$

to reduce the size of the argument of \sin . (For purposes of this exercise an angle is considered "sufficiently small" if its magnitude is not greater than 0.1 radians.) These ideas are incorporated in the following procedures:

```
(define (cube x) (* x x x))
```



```

(define (p x)
  ;; (Modified to show calls to p)
  (printf "(p ~a)~n" x)
  (- (* 3 x) (* 4 (cube x))))

(define (sine angle)
  (if (not (> (abs angle) 0.1))
      angle
      (p (sine (/ angle 3.0)))))

```

3.2.1 a.

How many times is the procedure `p` applied when `(sine 12.15)` is evaluated?

```

(sine 12.15)

(p 0.049999999999999996)
(p 0.1495)
(p 0.4351345505)
(p 0.9758465331678772)
(p -0.7895631144708228)

```

5 calls to `p`.

3.2.2 b.

What is the order of growth in space and number of steps (as a function of `a`) used by the process generated by the `sine` procedure when `(sine a)` is evaluated?

Logarithms answer, more or less, the question "how many times can I divide one number by another?" The second number is the *base*. So, consider log base 2 of 8: $8/2=4$, $4/2=2$, $2/2=1$; thus, $\log_2 8 = 3$.

The actual definition is that the log of a number is the exponent to which the base must be raised to equal that number. Thus, since $2^3 = 8$, then $\log_2 8 = 3$.

As can be seen by the single call to `p` in the body of `sine`, each recursive call reduces `angle` by a factor of 3; thus, `p` is going to be of $\theta(\log n)$. (The specific logarithmic base is effectively a constant, so all logarithmic processes are considered to be of the same order of computational complexity.)

4 Section 1.2.4 Exponentiation

4.1 Exercise 1.16: Iterative exponentiation in logarithmic time

4.1.1 Problem

Design a procedure that evolves an iterative exponentiation process that uses successive squaring and uses a logarithmic number of steps, as does **fast-expt**. (Hint: Using the observation that $(b^{n/2})^2 = (b^2)^{n/2}$, keep, along with the exponent n and the base b , an additional state variable a , and define the state transformation in such a way that the product ab^n is unchanged from state to state. At the beginning of the process a is taken to be 1, and the answer is given by the value of a at the end of the process. In general, the technique of defining an "invariant quantity" that remains unchanged from state to state is a powerful way to think about the design of iterative algorithms.)

4.1.2 Answer

Here's the code from the section for the original **fast-expt** algorithm.

```
(define (square n)
  (* n n))

(define (even? n)
  (= (remainder n 2) 0))

(define (fast-expt b n)
  (cond ((= n 0) 1)
        ((even? n) (square (fast-expt b (/ n 2))))
        (else (* b (fast-expt b (- n 1))))))
```

Note that the **even?** case in the **cond** is building up calls to **square** and the **else** is building up calls to *****. The stack of calls to **fast-expt** keeps building up until it bottoms out with the first case, after which all of the pending computations can be rolled back up.

This problem is to switch from this from a logarithmic recursive process to a logarithmic iterative process.

```
(define (fast-expt-2 b n)
  (define (fast-expt-iter b n a)
```

```

(cond ((= n 0) a)
      ((even? n) (fast-expt-iter (square b) (/ n 2) a))
      (else (fast-expt-iter b (- n 1) (* a b))))
(fast-expt-iter b n 1))

```

To evaluate this, we'll reuse the `compare` function used for Problem 1.11. To enable this, we wrap the calls to the text's `fast-expt` and the new `fast-expt-2` to take a single argument.

```

(define (expt-by-two n) (fast-expt 2 n))
(define (expt-by-two-2 n) (fast-expt-2 2 n))
(compare expt-by-two expt-by-two-2 20)

```

```

1: 2 (0 ms) 2 (0 ms) => #t
2: 4 (0 ms) 4 (0 ms) => #t
3: 8 (0 ms) 8 (1 ms) => #t
4: 16 (0 ms) 16 (0 ms) => #t
5: 32 (0 ms) 32 (0 ms) => #t
6: 64 (0 ms) 64 (0 ms) => #t
7: 128 (0 ms) 128 (0 ms) => #t
8: 256 (0 ms) 256 (0 ms) => #t
9: 512 (0 ms) 512 (0 ms) => #t
10: 1024 (0 ms) 1024 (0 ms) => #t
11: 2048 (0 ms) 2048 (0 ms) => #t
12: 4096 (0 ms) 4096 (0 ms) => #t
13: 8192 (0 ms) 8192 (0 ms) => #t
14: 16384 (0 ms) 16384 (0 ms) => #t
15: 32768 (0 ms) 32768 (0 ms) => #t
16: 65536 (0 ms) 65536 (0 ms) => #t
17: 131072 (0 ms) 131072 (0 ms) => #t
18: 262144 (0 ms) 262144 (0 ms) => #t
19: 524288 (0 ms) 524288 (0 ms) => #t
20: 1048576 (0 ms) 1048576 (0 ms) => #t

```

As can be seen, both are extremely fast and return identical results.

4.2 Exercise 1.17: Recursive integer multiplication with square and halve

4.2.1 Problem

The exponentiation algorithms in this section are based on performing exponentiation by means of repeated multiplication. In a similar way, one can

perform integer multiplication by means of repeated addition. The following multiplication procedure (in which it is assumed that our language can only add, not multiply) is analogous to the `expt` procedure:

```
(define (mult a b)
  (if (= b 0)
      0
      (+ a (* a (- b 1))))))
```

This algorithm takes a number of steps that is linear in b . Now suppose we include, together with addition, operations `double`, which doubles an integer, and `halve`, which divides an (even) integer by 2. Using these, design a multiplication procedure analogous to ‘fast-expt’ that uses a logarithmic number of steps.

4.2.2 Answer

This is a straightforward translation of the `fast-expt` code from the text to the multiplication problem. . . the problem (and solution) have exactly the same shape.

```
(define (double n) (* n 2))

(define (halve n) (/ n 2))

(define (even? n)
  (= (remainder n 2) 0))

(define (fast-mult a b)
  (cond ((= b 0) 0)
        ((even? b) (double (fast-mult a (halve b))))
        (else (+ a (fast-mult a (- b 1))))))
```

4.3 Exercise 1.18: Iterative integer multiplication with square and halve

4.3.1 Problem

Using the results of Exercise 1.16 and Exercise 1.17, devise a procedure that generates an iterative process for multiplying two integers in terms of adding, doubling, and halving and uses a logarithmic number of steps.

4.3.2 Answer

This solution is also straightforward. The only trick part is keeping straight what needs to be added and subtracted, and from where.

```
(define (fast-mult-2 a b)
  (define (fast-mult-iter a b c)
    (cond ((= b 0) c)
          ((even? b) (fast-mult-iter (double a) (half b) c))
          (else (fast-mult-iter a (- b 1) (+ c a)))))
  (fast-mult-iter a b 0))
```

4.4 Exercise 1.19: Logarithmic Fibonacci calculations

4.4.1 Problem

There is a clever algorithm for computing the Fibonacci numbers in a logarithmic number of steps. Recall the transformation of the state variables **a** and **b** in the **fib-iter** process of section 1.2.2: $a \leftarrow a + b$ and $b \leftarrow a$. Call this transformation T , and observe that applying T over and over again n times, starting with 1 and 0, produces the pair $\text{Fib}(n + 1)$ and $\text{Fib}(n)$. In other words, the Fibonacci numbers are produced by applying T^n , the n th power of the transformation T , starting with the pair $(1, 0)$. Now consider T to be the special case of $p = 0$ and $q = 1$ in a family of transformations T_{pq} , where T_{pq} transforms the pair (a, b) according to $a \leftarrow bq + aq + ap$ and $b \leftarrow bp + aq$. Show that if we apply such a transformation T_{pq} twice, the effect is the same as using a single transformation $T_{p'q'}$ of the same form, and compute p' and q' in terms of p and q . This gives us an explicit way to square these transformations, and thus we can compute T^n using successive squaring, as in the **fast-expt** procedure. Put this all together to complete the following procedure, which runs in a logarithmic number of steps:

```
(define (fib n)
  (fib-iter 1 0 0 1 n))

(define (fib-iter a b p q count)
  (cond ((= count 0) b)
        ((even? count)
         (fib-iter a
                   b
                   <??> ; compute p'
```

```

      <??>      ; compute q'
      (/ count 2)))
  (else (fib-iter (+ (* b q) (* a q) (* a p))
                  (+ (* b p) (* a q))
                  p
                  q
                  (- count 1)))))

```

4.4.2 Answer

Deriving the new transformation is slightly painful, and requires a lot of algebraic manipulation. The general procedure is to start with the initial transformation, apply it twice to the same inputs, and keep simplifying and reorganizing until you reach a form which has the same structure as the transformation, which will give the new p' and q' .

The form of a single transformation is

$$T_{pq}(a, b) = (bq + aq + ap, bp + aq)$$

Applying this twice gives

$$T_{pq}(T_{pq}(a, b)) = ((bp + aq)q + (bq + aq + ap)q + bq + aq + ap)p, \\ (bp + aq)p + (bq + aq + ap)q)$$

To make line lengths a bit shorter, we can examine the each part of the tuple separately; we'll call them a' and b' . To start, we'll by simplifying a' :

$$\begin{aligned} a' &= (bp + aq)q + (bq + aq + ap)q + bq + aq + ap)p \\ &= bpq + aq^2 + bq^2 + aq^2 + apq + bpq + apq + ap^2 \\ &= 2aq^2 + 2apq + 2bpq + bq^2 + bq^2 + ap^2 \end{aligned}$$

From here, we want to move things around into a form that is "shaped like" $bq + aq + ap$. To do this, we can note that no terms have both as and bs , so we can pull those terms out:

$$= b(2pq + q^2) + a(2q^2 + 2pq + p^2)$$

This gives us our q' , since q' is what should be multiplied by b in the new expression:

$$q' = 2pq + q^2$$

This lets us split up the other expression into its p' and q' components:

$$\begin{aligned} a' &= b(2pq + q^2) + a(q^2 + q^2 + 2pq + p^2) \\ &= b(2pq + q^2) + a(2pq + q^2) + a(q^2 + p^2) \end{aligned}$$

Sure enough, this gives us a value for p' , as well:

$$p' = q^2 + p^2$$

We can double-check this by making sure we get the same results for b' :

$$\begin{aligned} b' &= (bp + aq)p + (bq + aqap)q \\ &= bp^2 + apq + bq^2 + aq^2 + apq \\ &= aq^2 + 2apq + bq^2 + bp^2 \\ &= a(q^2 + 2pq) + b(q^2 + p^2) \end{aligned}$$

This matches up exactly. Thus, the double-application of T_{pq} is identical to $T_{p'q'}$, where $p' = q^2 + p^2$ and $q' = 2pq + q^2$. With that done, we just need to plug these formulae into the template provided by the exercise.

```
(define (fib-t n)
  (fib-iter 1 0 0 1 n))

(define (fib-t-iter a b p q count)
  (cond ((= count 0) b)
        ((even? count)
         (fib-t-iter a
                     b
                     (+ (square q) (square p))
                     (+ (* 2 p q) (square q))
                     (/ count 2)))
        (else (fib-t-iter (+ (* b q) (* a q) (* a p))
                           (+ (* b p) (* a q))
                           p
                           q
                           (- count 1))))))
```

5 Section 1.2.5 Greatest Common Divisors

5.1 TODO Exercise 1.20: Revisiting applicative order and normal order

5.1.1 Problem

The process that a procedure generates is of course dependent on the rules used by the interpreter. As an example, consider the iterative `gcd` procedure given above. Suppose we were to interpret this procedure using normal-order evaluation, as discussed in section 1.1.5. (The normal-order-evaluation rule for `if` is described in Exercise 1.5.) Using the substitution method (for normal order), illustrate the process generated in evaluating `(gcd 206 40)` and indicate the `remainder` operations that are actually performed. How many `remainder` operations are actually performed in the normal-order evaluation of `(gcd 206 40)`? In the applicative-order evaluation?

5.1.2 Answer

Recall that for applicative order, arguments are first evaluated, then the procedure is applied. For normal order, everything is fully expanded before the arguments are evaluated.

6 Section 1.2.6 Testing for Primality

6.1 Exercise 1.21: Using smallest-divisor

6.1.1 Problem

Use the `smallest-divisor` procedure to find the smallest divisor of each of the following numbers: 199, 1999, 19999.

6.1.2 Answer

First, the relevant code from section 1.2.6:

```
(require (planet neil/sicp:1:17))

(define (square n) (* n n))

(define (smallest-divisor n)
  (find-divisor n 2))
```



```
(define (find-divisor n test-divisor)
  (cond ((> (square test-divisor) n) n)
        ((divides? test-divisor n) test-divisor)
        (else (find-divisor n (+ test-divisor 1))))))
```

```
(define (divides? a b)
  (= (remainder b a) 0))
```

```
(define (prime? n)
  (= n (smallest-divisor n)))
```

```
(smallest-divisor 199)
```

An error occurred.

```
(smallest-divisor 1999)
```

An error occurred.

```
(smallest-divisor 19999)
```

An error occurred.

6.2 Exercise 1.22: Measuring runtime

6.2.1 Problem

Most Lisp implementations include a primitive called ‘runtime’ that returns an integer that specifies the amount of time the system has been running (measured, for example, in microseconds). The following ‘timed-prime-test’ procedure, when called with an integer *n*, prints *n* and checks to see if *n* is prime. If *n* is prime, the procedure prints three asterisks followed by the amount of time used in performing the test.

```
;; This code has been tweaked slightly to return true/false so
;; the return value can be used in tests. Also, only displays
;; output for prime numbers.
```

```
(define (timed-prime-test n)
  ;(display n)
  ;(display " ")
```

```

(start-prime-test n (runtime)))

(define (start-prime-test n start-time)
  (if (prime? n)
      (report-prime n (- (runtime) start-time))
      false))

(define (report-prime prime elapsed-time)
  (display prime)
  (display " *** ")
  (display elapsed-time)
  (newline)
  true)

```

Using this procedure, write a procedure ‘search-for-primes’ that checks the primality of consecutive odd integers in a specified range. Use your procedure to find the three smallest primes larger than 1000; larger than 10,000; larger than 100,000; larger than 1,000,000. Note the time needed to test each prime. Since the testing algorithm has order of growth of $\Theta(\sqrt{n})$, you should expect that testing for primes around 10,000 should take about $\sqrt{10}$ times as long as testing for primes around 1000. Do your timing data bear this out? How well do the data for 100,000 and 1,000,000 support the \sqrt{n} prediction? Is your result compatible with the notion that programs on your machine run in time proportional to the number of steps required for the computation?

6.2.2 Answer

First, a procedure to scan a range of consecutive odd numbers for primality:

```

(define (odd? n) (= (remainder n 2) 1))

;; Find primes in range from a to b
(define (search-for-primes a b)
  (if (< a b)
      (cond ((odd? a)
              (timed-prime-test a)
              (search-for-primes (+ a 2) b))
            (else
             (search-for-primes (+ a 1) b))))))

```

```
;; Find the first k primes larger than n
(define (find-k-primes k n)
  (if (odd? n)
      (if (> k 0)
          (if (timed-prime-test n)
              (find-k-primes (- k 1) (+ n 2))
              (find-k-primes k (+ n 2))))
      (find-k-primes k (+ n 1))))

;; Starting with =, find the first k higher primes;
;; then multiply n by 10 and repeat intervals times.
(define (prime-scan k intervals n)
  (find-k-primes k n)
  (if (> intervals 1) (prime-scan k (- intervals 1) (* n 10))))
```

Using this, finding the first three primes larger than 1,000 is easy: 1,009, 1,013, and 1,019.

```
(find-k-primes 3 1000)
```

An error occurred.

And for 10,000, 100,000, and 1,000,000::

```
(prime-scan 3 3 (expt 10 4))
```

An error occurred.

On my machine, calculating primality using this method for numbers around 10^4 takes about 8-9 microseconds, 10^5 takes about 25 microseconds, and 10^6 takes around 78 microseconds.

```
(display (* 8 (sqrt 10)))
(newline)
(display (* 25 (sqrt 10)))
```

```
25.298221281347036
```

```
79.05694150420949
```

These results match almost perfectly to the predicted execution time. The tweaked version of `prime-scan` makes it easy to test this at a broad range of magnitudes:

```
(prime-scan 1 13 (expt 10 4))
```

For a quick-and-dirty evaluation of this output, we'll munge it fast in the shell.

```
# Separate code block here so we can reuse these results easily in the
# next exercise
cut -f3 -d' '
```

And a quick script to churn through and verify the results against the predicted time:

```
echo Measured Predicted
for t in $(cut -f3 -d' '); do
  if [ -n "${last}" ]; then
    # dc is an ancient RPN calculator
    # space pushes a number, 'v' is sqrt
    # and 'p' prints the value on the top of the stack
    guess=$(dc -e"${last} 10v*p")
  fi
  echo $t $guess
  last=$t
done
```

Measured	Predicted
9	
25	27
79	75
278	237
804	834
2717	2412
8151	8151
15956	24453
50409	47868
159217	151227
507144	477651
1588455	1521432
4938960	4765365

These results continue to stay close to the predicted values which supports the analysis that run time is proportional to the number of steps required for the computation.

6.3 Exercise 1.23: Speeding up smallest-divisor

6.3.1 Problem

The ‘smallest-divisor’ procedure shown at the start of this section does lots of needless testing: After it checks to see if the number is divisible by 2 there is no point in checking to see if it is divisible by any larger even numbers. This suggests that the values used for ‘test-divisor’ should not be 2, 3, 4, 5, 6, ..., but rather 2, 3, 5, 7, 9, To implement this change, define a procedure ‘next’ that returns 3 if its input is equal to 2 and otherwise returns its input plus 2. Modify the ‘smallest-divisor’ procedure to use ‘(next test-divisor)’ instead of ‘(+ test-divisor 1)’. With ‘timed-prime-test’ incorporating this modified version of ‘smallest-divisor’, run the test for each of the 12 primes found in *Note Exercise 1-22::. Since this modification halves the number of test steps, you should expect it to run about twice as fast. Is this expectation confirmed? If not, what is the observed ratio of the speeds of the two algorithms, and how do you explain the fact that it is different from 2?

6.3.2 Answer

This is a little messy since we haven’t had higher-order functions introduced, yet, so here are all of the relevant functions rewritten to use the new `better-smallest-divisor` procedure.

```
(define (next test-divisor)
  (if (= test-divisor 2)
      3
      (+ test-divisor 2)))

(define (better-smallest-divisor n)
  (better-find-divisor n 2))

(define (better-find-divisor n test-divisor)
  (cond ((> (square test-divisor) n) n)
        ((divides? test-divisor n) test-divisor)
        (else (better-find-divisor n (next test-divisor)))))

(define (better-prime? n)
  (= n (better-smallest-divisor n)))
```

```

(define (better-timed-prime-test n)
  (display n)
  (display " ")
  (better-start-prime-test n (runtime)))

(define (better-start-prime-test n start-time)
  (if (better-prime? n)
      (report-prime n (- (runtime) start-time))
      false))

(define (better-find-k-primes k n)
  (if (odd? n)
      (if (> k 0)
          (if (better-timed-prime-test n)
              (better-find-k-primes (- k 1) (+ n 2))
              (better-find-k-primes k (+ n 2))))
      (better-find-k-primes k (+ n 1))))

;; Starting with =, find the first k higher primes;
;; then multiply n by 10 and repeat intervals times.

(define (better-prime-scan k intervals n)
  (better-find-k-primes k n)
  (if (> intervals 1) (better-prime-scan k (- intervals 1) (* n 10))))

```

Here are the 12 primes that are specified in the exercise:

```
(better-prime-scan 3 3 (expt 10 4))
```

An error occurred.

These results are very close to those for the original version... but, since these magnitudes are quite small relative to numbers that would have been expensive to calculate in 1996 (when SICP 2ed was published), it's difficult to differentiate. More useful is comparing at larger magnitudes:

```
(better-prime-scan 1 13 (expt 10 4))
```

```

10007 *** 5
100003 *** 15
1000003 *** 49

```

```

10000019 *** 151
100000007 *** 479
1000000007 *** 1519
10000000019 *** 4891
100000000003 *** 14766
1000000000039 *** 27112
10000000000037 *** 85590
100000000000031 *** 271503
1000000000000037 *** 855910
10000000000000061 *** 2741465

```

```
cut -f3 -d' '
```

```
(mapcar* 'append first second)
```

The first column is the original figures, and the second is the `better-*` version (it's a bit fussy to get the headers added into an `org-babel` block that combines two sets of output). These results at larger scales make the improved running time obvious.

6.4 Exercise 1.24: Putting the Fermat method to work

6.4.1 Problem

Modify the `timed-prime-test` procedure of Exercise 1.22 to use `fast-prime?` (the Fermat method), and test each of the 12 primes you found in that exercise. Since the Fermat test has $\theta(\log n)$ growth, how would you expect the time to test primes near 1,000,000 to compare with the time needed to test primes near 1000? Do your data bear this out? Can you explain any discrepancy you find?

6.4.2 Answer

First, the code from Section 1.2.6. Since we're pushing the input size larger than 4294967087, we can't use Racket's built-in `random`, so an external library from Planet (Racket's package repository) is used that does not cap the range (`williams/science/random-source`).

```
(require (planet williams/science/random-source))
```

```
(define (expmod base exp m)
  (cond ((= exp 0) 1)
```

```

      ((even? exp)
       (remainder (square (expmod base (/ exp 2) m))
                  m))
      (else
       (remainder (* base (expmod base (- exp 1) m))
                  m))))

(define (fermat-test n)
  (define (try-it a)
    (= (expmod a n n) a))
  (try-it (+ 1 (random-integer (- n 1)))))

(define (fast-prime? n times)
  (cond ((= times 0) true)
        ((fermat-test n) (fast-prime? n (- times 1)))
        (else false)))

```

Now, we need to modify the relevant procedures to use **fast-prime?**.

```

(define (fast-timed-prime-test n times)
  (fast-start-prime-test n times (runtime)))

(define (fast-start-prime-test n times start-time)
  (if (fast-prime? n times)
      (report-prime n (- (runtime) start-time))
      false))

(define (fast-find-k-primes k n times)
  (if (odd? n)
      (if (> k 0)
          (if (fast-timed-prime-test n times)
              (fast-find-k-primes (- k 1) (+ n 2) times)
              (fast-find-k-primes k (+ n 2) times)))
      (fast-find-k-primes k (+ n 1) times)))

(define (fast-prime-scan k intervals n times)
  (fast-find-k-primes k n times)
  (if (> intervals 1) (fast-prime-scan k (- intervals 1) (* n 10) times)))

```

Using 10 tests is super-fast:


```
(fast-prime-scan 1 28 (expt 10 3) 10)
```

An error occurred.

For a while, the increases in running time to appear to be more-or-less logarithmic: each power of ten increase in the input results in an increase to the running time of around 3-6 milliseconds.

There is some oddities, though: the jump in execution time between 10^9 and 10^{10} , which jumps up by almost an order of magnitude, and the jump between 10^{19} and 10^{20} , which multiplies the running time by around 25! A bit of thought reveals that a significant number occurs in the first range: 4294967087, which was the maximum integer that could be returned by Racket's `random`. Thus, we can guess that there may be a transparent promotion to a different internal representation that happens around that point, resulting in much higher constant costs to each operation. Presumably, a similar (and much for expensive) overflow happens at the higher threshold.

In the middle range, runtimes increase by around 20-60 milliseconds per power of ten increase to the inputs, which seems to confirm the prediction of logarithmic running time. In the upper range (10^{20} and higher), the numbers are a bit messier but the general trend of slow increase continues.

A simple improvement to the fidelity of these numbers could be made by testing multiple primes in each range, rather than just the first one found.

6.5 Exercise 1.25: A not-so-fast use of fast-expt

6.5.1 Problem

Alyssa P. Hacker complains that we went to a lot of extra work in writing `expmod`. After all, she says, since we already know how to compute exponentials, we could have simply written

```
(define (bad-expmod base exp m)
  (remainder (fast-expt base exp) m))
```

Is she correct? Would this procedure serve as well for our fast prime tester? Explain.

6.5.2 Answer

First, recall the relevant supporting code:

```

(define (square n)
  (* n n))

(define (even? n)
  (= (remainder n 2) 0))

(define (fast-expt b n)
  (cond ((= n 0) 1)
        ((even? n) (square (fast-expt b (/ n 2))))
        (else (* b (fast-expt b (- n 1))))))

```

And compare the problem's definition of `expmod` with the one used for Section 1.24:

```

(define (expmod base exp m)
  (cond ((= exp 0) 1)
        ((even? exp)
         (remainder (square (expmod base (/ exp 2) m))
                    m))
        (else
         (remainder (* base (expmod base (- exp 1) m))
                    m))))

```

The final result of both the original `expmod` and Alyssa's `bad-expmod` will be the same: they both calculate $base^{exp} \bmod m$. `bad-expmod` has to do a lot more work to achieve the same end, though, as it's manipulating much longer numbers: it generates the full exponential value before applying `remainder`... and `remainder` gets put through the wringer as it must divide that very large number by `m`. The original `expmod`, on the other hand, applies `remainder` at every step of the way, keeping the number in the range where it is both useful and easier to manipulate.

To test this, we'll set up a batch of procedures to use `bad-expmod`.

```

(define (bad-fermat-test n)
  (define (try-it a)
    (= (bad-expmod a n n) a))
  (try-it (+ 1 (random-integer (- n 1)))))

(define (bad-prime? n times)
  (cond ((= times 0) true)

```

```

        ((bad-fermat-test n) (bad-prime? n (- times 1)))
        (else false)))

(define (bad-timed-prime-test n times)
  (bad-start-prime-test n times (runtime)))

(define (bad-start-prime-test n times start-time)
  (if (bad-prime? n times)
      (report-prime n (- (runtime) start-time))
      false))

(define (bad-find-k-primes k n times)
  (if (odd? n)
      (if (> k 0)
          (if (bad-timed-prime-test n times)
              (bad-find-k-primes (- k 1) (+ n 2) times)
              (bad-find-k-primes k (+ n 2) times)))
      (bad-find-k-primes k (+ n 1) times)))

(define (bad-prime-scan k intervals n times)
  (bad-find-k-primes k n times)
  (if (> intervals 1) (bad-prime-scan k (- intervals 1) (* n 10) times)))

```

Here are some values for `fast-prime?` using the original `expmod`, using a small number of tests (just 10):

```
(fast-prime-scan 1 4 100 10)
```

An error occurred.

Barely any time at all, on the order of 10 microseconds. The new `bad-expmod` approach, however, does indeed live up to the name:

```
(bad-prime-scan 1 4 100 10)
```

An error occurred.

For even an input as small as 10^5 , the runtime is already nearing a second! This exercise is a great demonstration of potentially difficult to notice computation complexity bottlenecks, and the importance of picking the right algorithm for the job.

6.6 Exercise 1.26: A subtle slowdown in expmod

Louis Reasoner is having great difficulty doing Exercise 1-24. His `fast-prime?` test seems to run more slowly than his `prime?` test. Louis calls his friend Eva Lu Ator over to help. When they examine Louis's code, they find that he has rewritten the `expmod` procedure to use an explicit multiplication, rather than calling `square`:

```
(define (slow-expmod base exp m)
  (cond ((= exp 0) 1)
        ((even? exp)
         (remainder (* (slow-expmod base (/ exp 2) m)
                       (slow-expmod base (/ exp 2) m))
                    m))
        (else
         (remainder (* base (slow-expmod base (- exp 1) m))
                    m))))
```

"I don't see what difference that could make," says Louis. "I do." says Eva. "By writing the procedure like that, you have transformed the $\Theta(\log n)$ process into a $\Theta(n)$ process." Explain.

6.6.1 Answer

Once again, recall the original `expmod` procedure:

```
(define (expmod base exp m)
  (cond ((= exp 0) 1)
        ((even? exp)
         (remainder (square (expmod base (/ exp 2) m))
                    m))
        (else
         (remainder (* base (expmod base (- exp 1) m))
                    m))))
```

This is a lovely and subtle change. While appearing to be a simple in-place substitution of a procedure, it actually changes the single recursive call to `expmod` to be a tree of recursive calls, with two recursive calls at each internal node of the tree.

The original version divides the size of n by two at each stage...since n can only be divided by 2 at most $\log_2 n$ times, this gives the expected

complexity. While `slow-expmod` also divides the size of its argument by two, it also generates two recursive calls, one for each half. Thus, it does not reduce the size of the overall problem to be solved: while the tree has only $\log_2 n$ levels, there are 2^k subproblems at each level k . (Level 0 has a single problem; level 1 has $2^1 = 2$ problems. Each of those two problems generates two recursive children for $2^2 = 4$ problems at level two, and so on.

So, given $\Theta(\log_2 2^n)$, the log and the exponential cancel each other out (by the definition of logarithm, and the overall complexity is $\Theta(n)$).

6.7 Exercise 1.27: Fooling Fermat with Carmichael numbers

6.7.1 Problem

Demonstrate that the Carmichael numbers listed in Footnote 1.47 really do fool the Fermat test. That is, write a procedure that takes an integer n and tests whether a^n is congruent to $a \bmod n$ for every $a < n$, and try your procedure on the given Carmichael numbers.

6.7.2 Answer

```
(define (verify-fermat n)
  (define (verify-fermat-iter a n)
    (cond ((>= a n)
           true)
          ((= (expmod a n n) a)
           (verify-fermat-iter (+ a 1) n))
          (else
           false)))
  (verify-fermat-iter 1 n))

(define (descriptive-verify-fermat n)
  (display n)
  (if (verify-fermat n)
      (if (prime? n)
          (display ": prime and correctly passes the Fermat test")
          (display ": not prime and incorrectly passes the Fermat test"))
      (if (prime? n)
          (display ": prime and incorrectly fails the Fermat test")
          (display ": not prime and correctly fails the Fermat test")))
  (newline))
```

```
(define (fermat-scan-range a b)
  (descriptive-verify-fermat a)
  (if (< a b)
      (fermat-scan-range (+ a 1) b)))
```

This procedure does indeed show that the first six Carmichael numbers slip through the Fermat test.

```
(descriptive-verify-fermat 561)
(descriptive-verify-fermat 1105)
(descriptive-verify-fermat 1729)
(descriptive-verify-fermat 2465)
(descriptive-verify-fermat 2821)
(descriptive-verify-fermat 6601)
```

An error occurred.

```
(fermat-scan-range 1101 1109)
```

An error occurred.

For a bit more fun, we can turn this into a test for Carmichael numbers and find them ourselves. This could be much more fun with lists, `map`, and `filter`, but we haven't had them introduced, yet, so this sticks with printing out the relevant numbers.

```
(define (carmichael? n)
  (and (not (prime? n)) (verify-fermat n)))

(define (carmichael-scan-range a b)
  (if (carmichael? a) (printf "~a~n" a))
  (if (< a b) (carmichael-scan-range (+ a 1) b)))
```

Here's an example of using this to find all the Carmichael numbers under 10,000. As can be seen, the first six numbers mentioned in the text are all found using this method.

```
(carmichael-scan-range 1 100000)
```

6.8 Exercise 1.28: The Miller-Rabin test

6.8.1 Problem

One variant of the Fermat test that cannot be fooled is called the "Miller-Rabin test" (Miller 1976; Rabin 1980). This starts from an alternate form of Fermat's Little Theorem, which states that if n is a prime number and a is any positive integer less than n , then a raised to the $(n-1)$ st power is congruent to 1 modulo n . To test the primality of a number n by the Miller-Rabin test, we pick a random number $a < n$ and raise a to the $(n-1)$ st power modulo n using the `expmod` procedure. However, whenever we perform the squaring step in `expmod`, we check to see if we have discovered a "nontrivial square root of 1 modulo n ," that is, a number not equal to 1 or $n-1$ whose square is equal to 1 modulo n . It is possible to prove that if such a nontrivial square root of 1 exists, then n is not prime. It is also possible to prove that if n is an odd number that is not prime, then, for at least half the numbers $a < n$, computing $a^{(n-1)}$ in this way will reveal a nontrivial square root of 1 modulo n . (This is why the Miller-Rabin test cannot be fooled.) Modify the `expmod` procedure to signal if it discovers a nontrivial square root of 1, and use this to implement the Miller-Rabin test with a procedure analogous to `fermat-test`. Check your procedure by testing various known primes and non-primes. Hint: One convenient way to make `expmod` signal is to have it return 0.

6.8.2 Answer

This has some ugly bits...judicious use of `let` (which isn't introduced until the next section) would again simplify some of these expressions.

```
;; Test whether i is a nontrivial square root of 1 modulo m
(define (nontrivial-sqrt-mod? i m)
  (and (not (= i 1))
        (not (= i (- m 1)))
        (= (remainder (square i) m) 1)))

(define (mr-expmod base exp m)
  (define (maybe-continue i)
    (if (or (= i 0) (nontrivial-sqrt-mod? i m))
        0
        (remainder (square i) m)))
```

```

(cond ((= exp 0) 1)
      ((even? exp)
       (maybe-continue (mr-expmod base (/ exp 2) m)))
      (else
       (remainder (* base (mr-expmod base (- exp 1) m))
                   m))))

(define (mr-test n)
  (define (try-it a)
    ;; We don't need to check if the return of mr-expmod = 0,
    ;; as it is always the case that a>1.
    (= (mr-expmod a (- n 1) n) 1))
  (try-it (+ 1 (random-integer (- n 1)))))

(define (mr-prime? n times)
  ;; We have to special-case n=1 and n=2.
  ;; (Note prime? incorrectly reports 1 as prime, and
  ;; fast-prime also fails outright.)
  (cond ((= times 0) true)
        ((= n 1) false)
        ((= n 2) true)
        ((mr-test n) (mr-prime? n (- times 1)))
        (else false)))

(define (mr-timed-prime-test n times)
  (mr-start-prime-test n times (runtime)))

(define (mr-start-prime-test n times start-time)
  (if (mr-prime? n times)
      (report-prime n (- (runtime) start-time))
      false))

(define (mr-find-k-primes k n times)
  (if (odd? n)
      (if (> k 0)
          (if (mr-timed-prime-test n times)
              (mr-find-k-primes (- k 1) (+ n 2) times)
              (mr-find-k-primes k (+ n 2) times)))
      (mr-find-k-primes k (+ n 1) times)))

```



```
(define (mr-prime-scan k intervals n times)
  (mr-find-k-primes k n times)
  (if (> intervals 1) (mr-prime-scan k (- intervals 1) (* n 10) times)))
```

Trying this out, `mr-prime?` successfully detects and discriminates primes and non-primes.

```
(mr-prime-scan 3 3 (expt 10 4) 3)
```

An error occurred.

When we test it against the Carmichael numbers found in the previous exercise, it correctly identifies them all as non-prime:

```
(define (mr-check n)
  (printf "~a: ~a~n" n (mr-prime? n 5)))
(mr-check 561)
(mr-check 1105)
(mr-check 1729)
(mr-check 2465)
(mr-check 2821)
(mr-check 6601)
(mr-check 8911)
(mr-check 10585)
(mr-check 15841)
(mr-check 29341)
(mr-check 41041)
(mr-check 46657)
(mr-check 52633)
(mr-check 62745)
(mr-check 63973)
(mr-check 75361)
```

An error occurred.

Additional, testing the same range of values as was examined in exercise 1.24 shows that the same set of primes is discovered.

```
(mr-prime-scan 1 28 (expt 10 3) 10)
```

An error occurred.